

# HPC the Easy Way

Tools and techniques for making the most of your resources

RSE Sheffield Seminar Series  
University of Sheffield  
30 July 2019

Phil Tooley  
HPC Application Analyst



Experts in numerical software and  
High Performance Computing

# Outline

## Common HPC Problems

Using the HPC more efficiently

The real world — ShARC

## HPC Package Managers

Conda

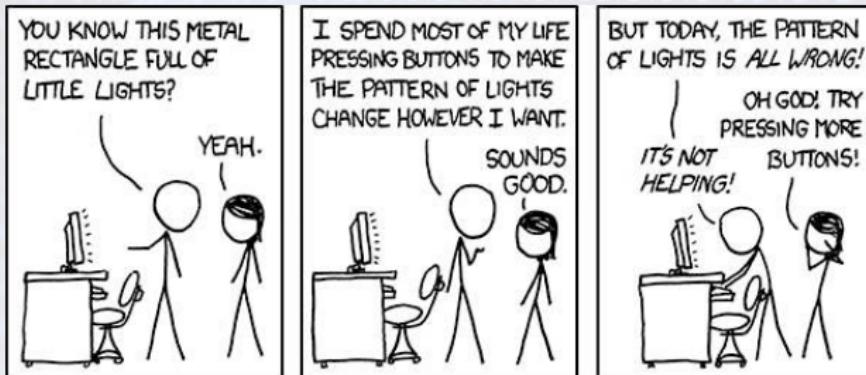
Spack

The POP-COE

# Common HPC Problems

## Two common HPC problems

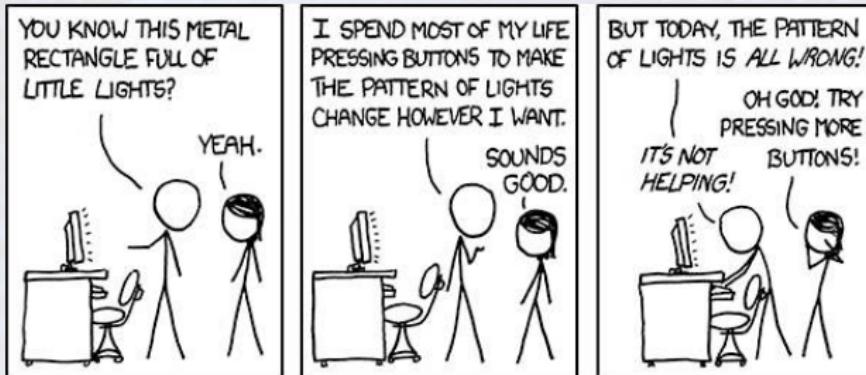
- ▶ Why is my job still queuing?
- ▶ How do I install <package>?



# Common HPC Problems

## Two common HPC problems

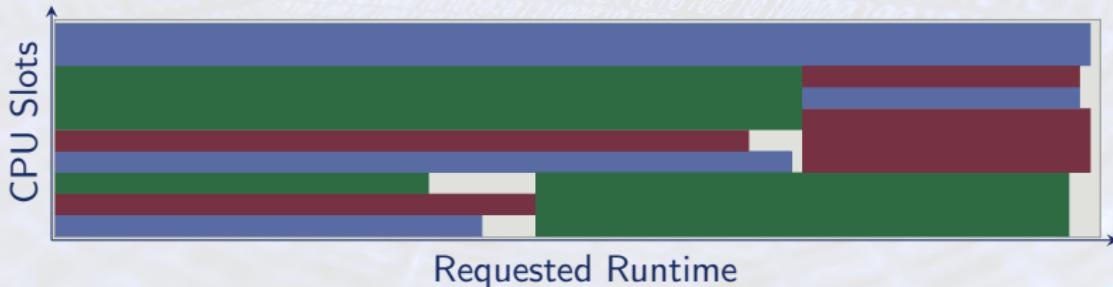
- ▶ Why is my job still queuing?
- ▶ How do I install <package>?



# What the Scheduler does

## A bin-packing problem

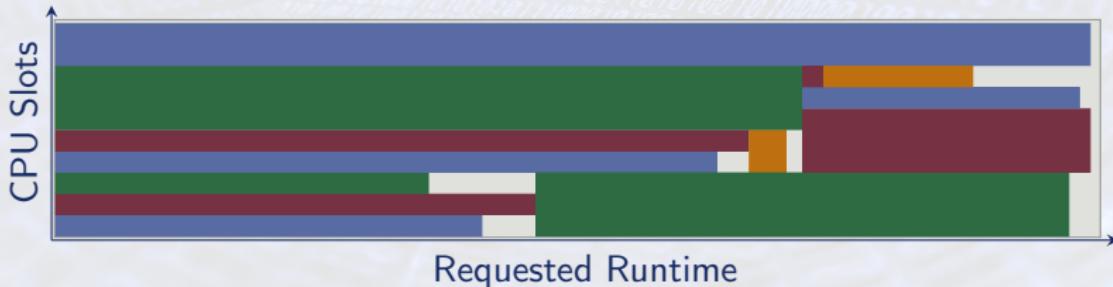
- ▶ Plans how to map jobs into nodes as efficiently as possible
- ▶ No job should wait "too long"
- ▶ Everyone should get a "fair share"
- ▶ Small jobs fill gaps around big ones



# What the Scheduler does

## A bin-packing problem

- ▶ Gaps appear as jobs finish early or are cancelled
- ▶ Scheduler backfills gaps as best it can
- ▶ Smaller jobs have more chances to backfill
- ▶ Ask for only what you actually need



# The real world picture - ShARC

## Mining the scheduler data

- ▶ Who is using ShARC?
- ▶ How are they using it?
- ▶ How efficiently are they using it?

# The real world picture - ShARC

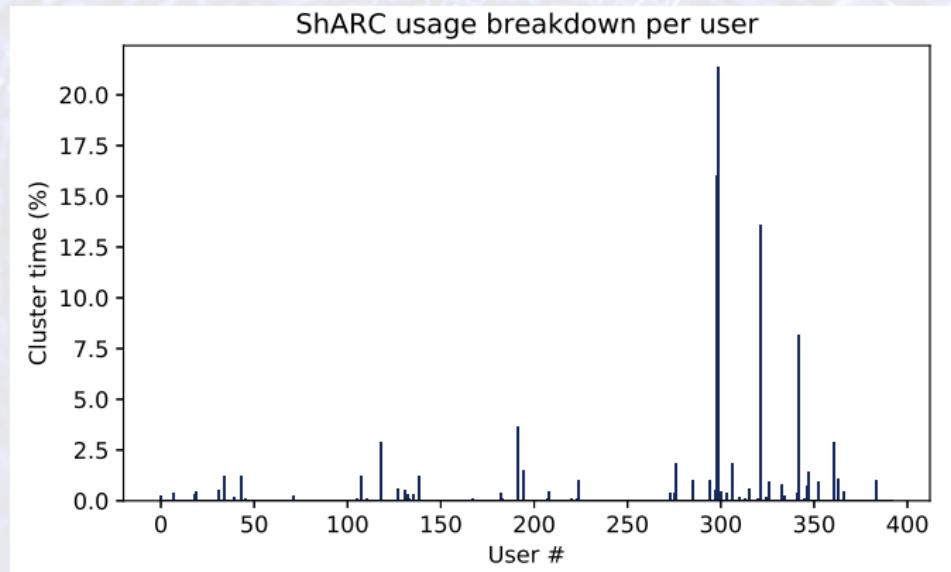
## Mining the scheduler data

- ▶ Who is using ShARC?
- ▶ How are they using it?
- ▶ How efficiently are they using it?

## The dataset

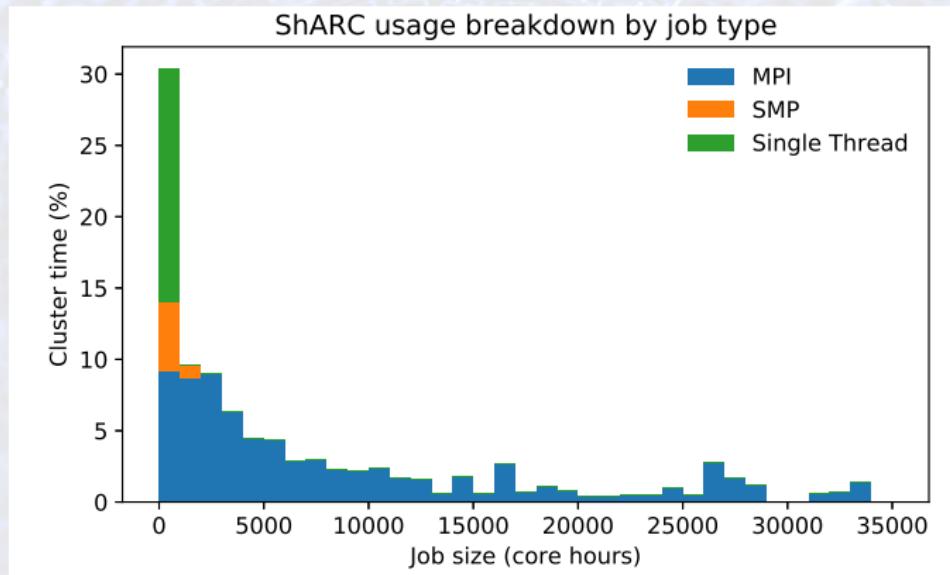
- ▶ Jobs started between 1/7/2017 – 30/6/2018
- ▶ Only public node data
- ▶ Failed jobs removed
- ▶ Sysadmin test jobs removed

# User Breakdown



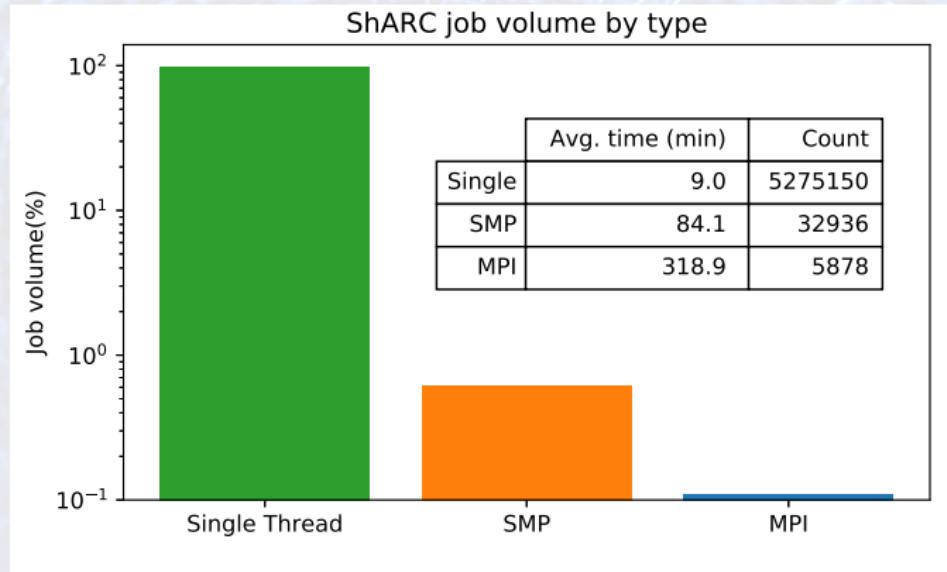
- ▶ 539 unique users
- ▶ Heaviest 3 users consumed over 50% of available cpu time

# Job breakdown



- ▶ Most time is spent running MPI jobs
- ▶ ~ 75% MPI vs. ~ 25% single node/thread

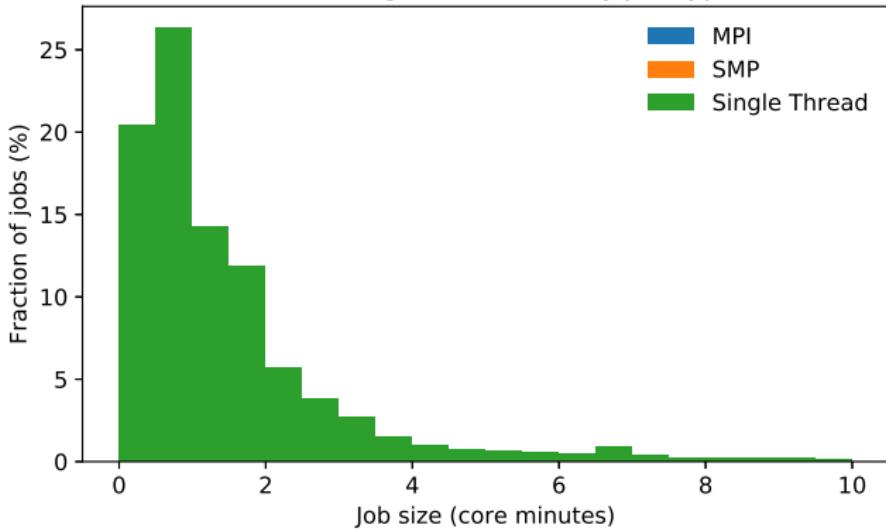
## Jobs breakdown



- ▶ Huge volume of very short jobs
- ▶ Heaviest users submitting  $> 10^6$  short jobs each!

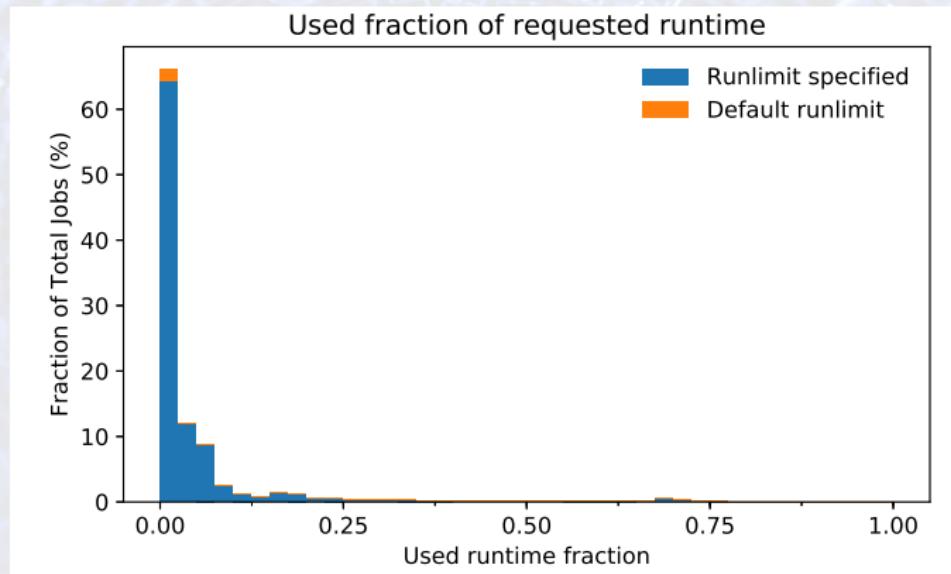
## Jobs breakdown

ShARC usage breakdown by job type



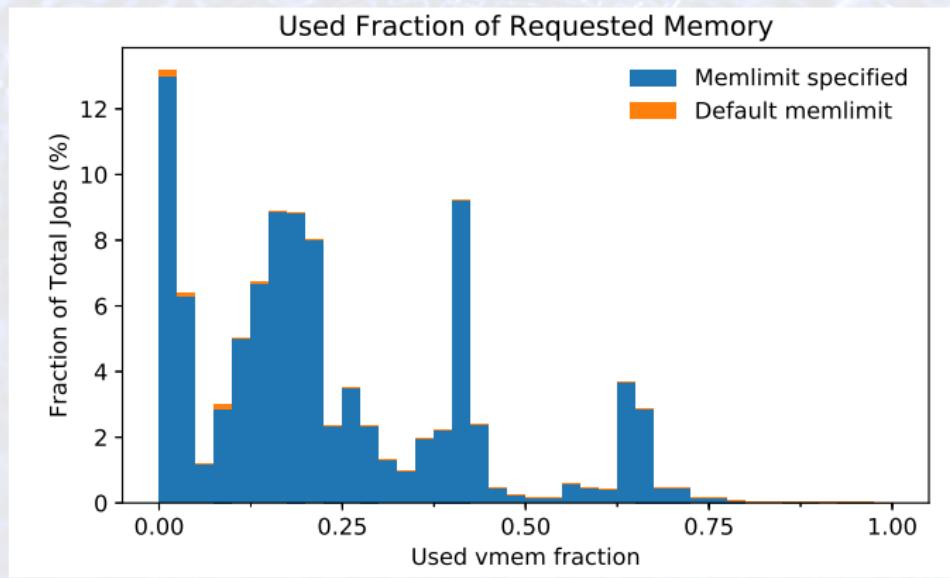
- ▶ ~ 50% of ShARC jobs shorter than 1 minute
- ▶ 50% of scheduler effort spent on only 0.4% of cpu time!

# Runtime Requests and Usage



- ▶ Most over-request walltime by at least an order of magnitude
- ▶ → Lots of missed opportunities to backfill gaps!

# Memory Requests and Usage



- ▶ Majority of users explicitly request memory
- ▶ Better usage, but still lots of over-requesting

# Getting Feedback from the Scheduler

## Accounting Information

- ▶ ShARC/Iceberg

```
$ qacct -j $jobid
```

- ▶ Bessemer

```
$ sacct -j $jobid
```

- ▶ Records basic performance information about job

- Requested resources (time, memory etc.)
- Actual runtime
- Actual memory usage
- Useful CPU time

# Accounting Information

```
qacct -j 1150879
```

qname	all.q
hostname	sharc-node147.shef.ac.uk
owner	ac1mpt
job_number	1150879
submission_time	2018-04-16 10:00:43
start_time	2018-04-16 10:00:54
end_time	2018-04-19 10:34:48
exit_status	0
ru_wallclock	261234
granted_pe	mpi
slots	220
cpu	57314572.128644
category	-u ac1mpt -l h_rt=345600,h_vmem=2G -pe mpi 220 -P SHEFFIELD
maxvmem	150.63G

# Resource Rules of Thumb

## Runtime

- ▶ Check `ru_wallclock` — actual run time
- ▶ Request 1.5–2× `ru_wallclock`

# Resource Rules of Thumb

## Runtime

- ▶ Check `ru_wallclock` — actual run time
- ▶ Request 1.5–2× `ru_wallclock`

## Memory

- ▶ Check `maxvmem` — peak job memory usage
- ▶ Request 1.5–2× `maxvmem`
- ▶ Remember requests are *per core*

# Resource Rules of Thumb

## Runtime

- ▶ Check `ru_wallclock` — actual run time
- ▶ Request 1.5–2× `ru_wallclock`

## Memory

- ▶ Check `maxvmem` — peak job memory usage
- ▶ Request 1.5–2× `maxvmem`
- ▶ Remember requests are *per core*

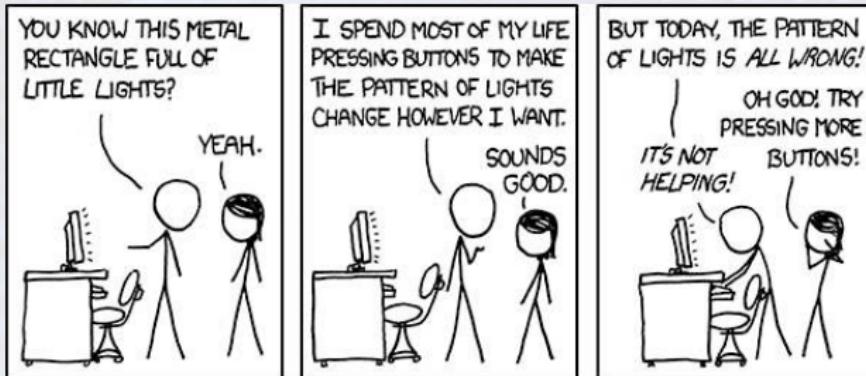
## Efficiency

- ▶ Check `cpu` — actual cpu usage
- ▶ Ensure  $\text{cpu} \simeq \text{ru\_wallclock} \times \text{slots}$

# Common HPC Problems

## Two common HPC problems

- ▶ Why is my job still queuing?
- ▶ How do I install <package>?



# Automating Software Installation

## Package Managers

- ▶ Automate installation/removal of software
- ▶ Manage installation of required dependencies
- ▶ Curate package repositories
- ▶ Document and reproduce environments

Focus on just two:

CONDA



Spack

# Conda

Pre-built packages for Python, R, etc.

- ▶ Originally for Anaconda Python distribution
- ▶ Microsoft provided R packages
- ▶ Low level numerical support libraries
- ▶ Intel Python with MKL optimised Numpy/Scipy
- ▶ Designed for users to install what they need



# Installing Conda

## Personal machine — Windows, Mac, Linux

- ▶ Two versions:
- ▶ Anaconda — Full distribution with hundreds of packages
- ▶ Miniconda — Just Conda and Python
- ▶ Download from [anaconda.com](http://anaconda.com) and run installer

# Installing Conda

## Personal machine — Windows, Mac, Linux

- ▶ Two versions:
- ▶ Anaconda — Full distribution with hundreds of packages
- ▶ Miniconda — Just Conda and Python
- ▶ Download from [anaconda.com](http://anaconda.com) and run installer

## ShARC, Bessemer, Iceberg

- ▶ Already installed:

```
$ module load conda
```

# Installing and Managing Packages

## Conda Environments

- ▶ Collections of packages and their dependencies
- ▶ Isolate individual projects
- ▶ Test/use multiple versions of a package
- ▶ Easily capture and reproduce environment elsewhere

# Installing and Managing Packages

## Conda Environments

- ▶ Collections of packages and their dependencies
- ▶ Isolate individual projects
- ▶ Test/use multiple versions of a package
- ▶ Easily capture and reproduce environment elsewhere

## Creating Environments

```
$ conda create --name myenv numpy pystan  
$ source activate myenv
```

# Installing and Managing Packages

## Lots of customization options

- ▶ Choose Python version:

```
$ conda create --name myenv numpy pystan python=3.7
```

# Installing and Managing Packages

## Lots of customization options

- ▶ Choose Python version:

```
$ conda create --name myenv numpy pystan python=3.7
```

- ▶ Package versions:

```
$ conda create --name myenv numpy pystan=2.17.1
```

# Installing and Managing Packages

## Lots of customization options

- ▶ Choose Python version:

```
$ conda create --name myenv numpy pystan python=3.7
```

- ▶ Package versions:

```
$ conda create --name myenv numpy pystan=2.17.1
```

- ▶ Other channels, e.g Intel Python

```
$ conda create --channel intel --name myenv numpy
```

# Installing and Managing Packages

## Lots of customization options

- ▶ Choose Python version:

```
$ conda create --name myenv numpy pystan python=3.7
```

- ▶ Package versions:

```
$ conda create --name myenv numpy pystan=2.17.1
```

- ▶ Other channels, e.g Intel Python

```
$ conda create --channel intel --name myenv numpy
```

- ▶ Non Python environments e.g R:

```
$ conda create --channel r --name myRenv r rstudio
```

# Using Environments

## Activating and deactivating

- ▶ “Activate” an environment to use it:

```
$ conda activate myenv
```

# Using Environments

## Activating and deactivating

- ▶ “Activate” an environment to use it:

```
$ conda activate myenv
```

- ▶ Installed Packages are now available to use:

```
$ python
Python 3.6.8 (default, Mar 10 2019, 17:04:16)
    >>> module load pystan
    >>> module load numpy
    >>> # etc...
```

# Using Environments

## Activating and deactivating

- ▶ “Activate” an environment to use it:

```
$ conda activate myenv
```

- ▶ Installed Packages are now available to use:

```
$ python
Python 3.6.8 (default, Mar 10 2019, 17:04:16)
    >>> module load pystan
    >>> module load numpy
    >>> # etc...
```

- ▶ “Deactivate” the environment to exit:

```
$ conda deactivate
```

# Using Environments

## Installing extra packages

- ▶ Can add extra packages to the environment

```
$ conda activate myenv  
$ conda install scipy scikit-learn #etc...
```

- ▶ And remove unneeded ones

```
$ conda remove scikit-learn #etc...
```

# Using Environments

## Installing extra packages

- ▶ Can add extra packages to the environment

```
$ conda activate myenv  
$ conda install scipy scikit-learn #etc...
```

- ▶ And remove unneeded ones

```
$ conda remove scikit-learn #etc...
```

## Updating packages

- ▶ Update all packages to the latest version:

```
$ conda activate myenv  
$ conda update --all
```

# Exporting Environments

## Preserving Environments

- ▶ Export complete list of packages with versions to a file:

```
$ conda env export --name myenv > myenv.txt
```

# Exporting Environments

## Preserving Environments

- ▶ Export complete list of packages with versions to a file:

```
$ conda env export --name myenv > myenv.txt
```

## Recreating Environments

- ▶ Now take that package list to another machine:

```
$ conda create --name myenv_clone -f myenv.txt
```

- ▶ `myenv_clone` is now an exact copy of `myenv`

- Collaboration with other users
  - Porting to new machines
  - Publishing for reproducibility

- ▶ Plain text file listing packages — can also be created/edited by hand

# Conda — Summary

## Python and R Package Management

- ▶ Designed for portability and reproducibility
- ▶ Rapidly install Python, R etc. packages
- ▶ Full control of package versioning
- ▶ Maintain multiple custom package environments
- ▶ Export, share and duplicate environments



## Build scientific packages from source

- ▶ Primarily designed for HPC package management
- ▶ Build optimised packages for specific system
- ▶ “Recipes” to install over 3000 packages
- ▶ Interoperates with already installed packages
- ▶ For sysadmins and end-users



# Installing Spack

## Requirements

- ▶ Python >= 2.6
- ▶ A working compiler (gcc, intel, pgi, etc.)

# Installing Spack

## Requirements

- ▶ Python >= 2.6
- ▶ A working compiler (gcc, intel, pgi, etc.)

## Installation

```
$ cd $HOME  
$ git clone https://github.com/spack/spack.git  
$ export SPACK_ROOT="$HOME/spack"  
$ source $SPACK_ROOT/share/spack/setup-env.sh
```

- ▶ Install as user in homedir
- ▶ Use .bashrc to automatically set up

# Configuring Spack

## Compiler autodetection

```
$ spack compilers
==> Available compilers
-- gcc sles12-x86_64 -----
gcc@4.8
```

# Configuring Spack

## Compiler autodetection

```
$ spack compilers
==> Available compilers
-- gcc sles12-x86_64 -----
gcc@4.8
```

## Additional compilers

```
$ module load gcc/8.1.0
$ spack compiler find
==> Added 1 new compiler:
    gcc@8.1.0
```

# Configuring Spack

## System packages

- ▶ Often want to use some system packages, e.g:
  - Vendor optimised MPI
  - System supplied BLAS/LAPACK
  - Avoid compiling again
- ▶ Specify in `packages.yaml`

```
# /home/phil/.spack/linux/packages.yaml
packages:
    netlib-lapack:
        modules: lapack/3.8.0
        buildable: False
```

# Installing Packages

## Search available packages

```
$ spack list mpi
==> 21 packages.
intel-mpi    mpibash    mpiblast    mpich    openmpi ...
```

# Installing Packages

## Search available packages

```
$ spack list mpi
==> 21 packages.
intel-mpi    mpibash    mpiblast    mpich    openmpi ...
```

## Install a package

- ▶ Install “preferred” version

```
$ spack install openmpi
```

- ▶ Specify a version

```
$ spack install openmpi@2.1.0
```

# Spack — Summary

## HPC Package Management

- ▶ A heavy duty package manager
- ▶ Designed for flexibility and control
- ▶ Integration with system modules and packages
- ▶ Full control of package versioning
- ▶ Build optimised packages from source





# Parallel Performance Optimization and Productivity

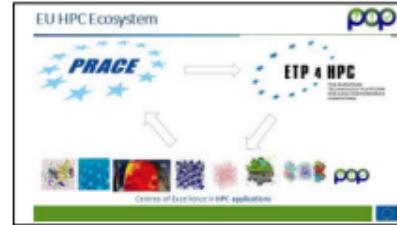


EU H2020 Centre of Excellence (CoE)

Grant Agreement No 824080

1 December 2018 – 30 November 2021

- A **Centre of Excellence**
  - On **Performance Optimisation and Productivity**
  - Promoting **best practices in parallel programming**
- Providing **FREE Services**
  - Precise understanding of application and system behaviour
  - Suggestion/support on how to refactor code in the most productive way
- **Horizontal**
  - Transversal across application areas, platforms, scales
- **For (EU) academic AND industrial codes and users !**



# Partners



- Who?

- BSC, ES (coordinator)
- HLRS, DE
- IT4I, CZ
- JSC, DE
- NAG, UK
- RWTH Aachen, IT Center, DE
- TERATEC, FR
- UVSQ, FR



JÜLICH  
SUPERCOMPUTING  
CENTRE



## A team with

- Excellence in performance tools and tuning
- Excellence in programming models and practices
- Research and development background AND proven commitment in application to real academic and industrial use cases



## Why?

- Complexity of machines and codes
  - ⇒ Frequent lack of quantified understanding of actual behaviour
  - ⇒ Not clear most productive direction of code refactoring
- Important to maximize efficiency (performance, power) of compute intensive applications and productivity of the development efforts

## What?

- Parallel programs, mainly MPI/OpenMP
  - Although also CUDA, OpenCL, OpenACC, Python, ...

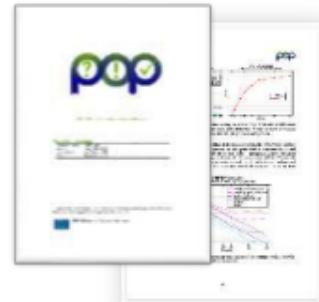


# FREE Services provided by the CoE



- Parallel Application Performance Assessment

- Primary service
- Identifies performance issues of customer code (at customer site)
- If needed, identifies the root causes of the issues found and qualifies and quantifies approaches to address them (recommendations)
- Combines former Performance Audit (?) and Plan (!)
- Medium effort (1-3 months)



- Proof-of-Concept (✓)

- Follow-up service
- Experiments and mock-up tests for customer codes
- Kernel extraction, parallelisation, mini-apps experiments to show effect of proposed optimisations
- Larger effort (3-6 months)

```
<!DOCTYPE html>
<html id="home-layout">
  <head>
    <meta http-equiv="content-type" content="text/html; charset=UTF-8"/>
    <title>Source Code Pro</title>
    <!-- made with <3 and AFKWD -->
    <meta name="keywords" content="sans, monospace, open source, coding, for developers" />
    <link rel="stylesheet" type="text/css" href="css/app.css"/>
  </head>
  <body>
    <div id="main">
```

Note: Effort shared between our experts and customer!



# The Process ...



## When?

December 2018 – November 2021

## How?

- Apply
  - Fill in small questionnaire describing application and needs  
<https://pop-coe.eu/request-service-form>
    - Questions? Ask [pop@bsc.es](mailto:pop@bsc.es)
  - Selection/assignment process
  - Install tools @ your production machine (local, PRACE, ...)
  - Interactively: Gather data → Analysis → Report

The screenshot shows the 'Request Service Form' page on the POP website. The left sidebar has links for Home, Blog, Newsletter, Partners, Tools, Services, and a highlighted 'Request Service Form'. The main form area has sections for 'Contact Details' (Applicant's Name, Institution, e-mail), 'Code' (Name of the code, description, date), 'Classification' (radio buttons for Core developer, Main developer, User), 'Access to resources' (radio buttons for 10%, 50%, 100%), 'Programming languages' (checkboxes for C/C++, Fortran, OpenMP, Python, Julia, R, others), 'Parallel programming models' (checkboxes for MPI, OpenMP, OpenMP, Thrust, CUDA, OpenCL, others), and 'Performance Service' (Service request dropdown set to 'Select', text area for problem description).

